



Offensive language exploratory analysis

Maša Kljun, Matija Teršek

Abstract

In this paper we focus on the exploratory analysis of 10 different subgroups of hate speech. We use natural language processing techniques in order to find the underlying structure and connections/relations between the subgroups. We focus on data extracted from Twitter and online forums. First we use classic approaches, such as TF-IDF, BoW, and LDA, then we move on to more sophisticated methods such as embeddings. [Note to professor: More to come in the next submission.](#)

Keywords

Hate speech, TF-IDF, embeddings, exploratory analysis, NLP ...

Advisors: Slavko Žitnik

Introduction

In the last few years social media grew exponentially and with it also the ability of people to express themselves online. By enabling people to write on different online platforms without even identifying themselves it lead to a new era of freedom of speech. As this new medium for communication and writing brought many positive things, it also has its downside. Social media has become a place where heated discussions happen and often result in insults and hatred. It is an important task to recognize hate speech and to prevent it.

Hate speech is defined as *abusive or threatening speech or writing that expresses prejudice against a particular group, especially on the basis of race, religion, or sexual orientation* [1]. We can see that the definition is very vague. Having said that, the goal of our paper is to help distinguish different types of hate speech and find the specific keywords of its subgroups in order to explain its structure. This could help with its identification and classification. In this paper we focus on ten subgroups of hate speech - *abusive, hateful, spam, general hate speech, profane, offensive, cyberbullying, racism, sexism, and benevolent sexism*. With understanding the structure of these groups, the goal is to also find similarities and connections between them.

There has been done a lot of research regarding the hate speech, however these works are usually focused on the classification of hate speech. One of the first works include [2] who built the decision tree based classifier Smokey for abusive message recognition and classification. Some other works that focus mainly on classification include [3] who compare the classification accuracy of models trained on expert and

amateur annotations, [4] who use convolutional neural networks for classification into four predefined categories, and [5] who use different natural language processing techniques for expanding datasets with emotional information for better classification. In the last years, especially deep learning models are often used for detection and classification of hate speech, such as [6] who propose a sophisticated method that is a combination of a deep neural network architecture with transfer learning. There is also a lot of related work that focuses on creating large datasets such as [7] who create a large-scale, multilingual, expert based dataset of hate speech.

What is less common in the research area of hate speech is analysis of relationships between different types of hate speech and the importance of specific keywords. Some examples include [8], who try to separate bullying from other social media posts and try to discover topic of bullying using topic modeling with Latent Dirichlet Allocation (LDA). [9] model hate speech against immigrants on Twitter in Spain. They try to find underlying topic of hate speech using LDA, discovering features of different dimensions of hatespeech, including foul language, humiliation, irony, etc. [10] conduct a survey about hate speech detection and describe key areas that have been explored, regarding the topic modeling, as well as sentiment analysis.

[Note to professor: Our goal for the future is first a deeper analysis of the obtained data. We will try to extract some underlying topics for each label and present the main findings \(LDA\). Use TF-IDF to find the most common \$n\$ -grams for each label. And in the future focus on the newer approaches, such as embedding with the further goal of exploring the](#)

relationships between types of hate speech. We hope that we can determine a more precise plan on the first submission defense.

This paper is organized as follows: TODO at the end of the report

Data

We use six publicly available datasets for our exploratory analysis. We combine datasets [3], [6], and [11] into one large dataset (referred to as Dataset SRB) as they include same categories of hate speech. We make labels *sexism*, *racism*, and *both* from [3] and [6]. The third dataset ([11]) that we use contains label *hostile sexism*, where marked tweets are already included in the first two datasets under *sexism*, and label *benevolent sexism*, which we rename to *benevolent*. We obtain a dataset with 6069 samples that are labeled either *sexism*, *racism*, *both*, or *benevolent*. The fourth dataset (referred to as Dataset AHS)[12] that we use has 3 categories - *abusive*, *hateful*, *spam*. As this is the original dataset no additional merging is needed. We obtain a dataset with 13776 tweets with the mentioned labels. Note that we exclude *None* label from both datasets, as we do not need it for the analysis. We show the distribution of individual categories from datasets SRB and AHS in Figures 1 and 2, respectively. Note that the numbers of samples might not match the numbers in the original papers, due to the Twitter removing the tweets, making them unavailable for us to analyze. We also provide an example for each label.

Racism - "He can't be a server at our restaurant, that beard makes him look like a terrorist." Everyone laughs. #fuck-thanksgiving

Sexism - #katieandnikki stop calling yourselves pretty and hot..you're not and saying it a million times doesn't make you either...STFU

Benevolent - It's "NEXT to every successful man, there's a woman"

Spam - RT @OnlyLookAtMino: [!!] #WINNER trending #1 on melon search

Abusive - You Worried About Somebody Bein Ugly... Bitch You Ugly...

Hateful - i hope leaders just kick retards that fake leave teams today

Additionally we use the dataset of comments extracted from the League of Legends community [13]. We preprocess the dataset given in the SQL format to a more readable CSV form and keep only the posts that are annotated as harassment. We obtain 259 examples of cyberbullying examples. The sixth dataset that we use was designed for the problem of the hate speech identification and classification, but we use the labels from the train and test set and merge them into one big dataset




Figure 1. Distribution of tweets in SRB dataset. This figure shows the distribution of hate speech categories in the SRB dataset. We can see that *sexism* and *benevolent* are well represented, whereas *racism* and *both* are far less frequent. Original set contains more tweets labeled *racism*, but due to their removal we cannot obtain them.




Figure 2. Distribution of tweets in AHS dataset. We see that the *spam* is the most represented label in the dataset, which represents the majority of the dataset. This is followed by the *abusive* tweets and there is the least *hateful* tweets. We can see that categories in this dataset are well represented.

that we use for our analysis. It provides tags of *hatespeech*, *profane*, and *offensive*, so we refer to the dataset as HPO. It consists of 2549 tweets, distribution of which can be seen in Figure 3. We again provide an example for each of the labels.

Cyberbullying - plot twist she's a fggt

Hatespeech - Johnson you liar. You don't give a flying one for the Irish

Offensive - #FuckTrump And retired porn star Melania too.

Profane - Fuck Trump and anybody who voted for that Lyin POS! #FuckTrump




Figure 3. Distribution of tweets in HPO dataset. The most used label is *hatespeech*. It is followed by *profane* and then *offensive*, which have a similar number of tweets.

We also use the dataset of Wikipedia comments [14], that




Figure 4. Distribution of tweets in TOITI. We see that most of the comments are labeled as *toxic*. Around half of them are *obscene* and around half are also labeled as *insult*. *Identity hate* and *threat* are far more uncommon in this dataset.

are marked as either *toxic*, *sever toxic*, *obscene*, *identity hate*, *threat*, and *insult*. We merge the first two categories into *toxic*. It is important to note that each comment in this dataset might have multiple labels, so the results for those tags might be similar. Original dataset contains 159571 tweets, 16225 of which are labeled. We show the distribution of the labels in Figure 4. We denote this dataset as TOITI in the future text.

Threat - SHUT UP, YOU FAT POOP, OR I WILL KICK YOUR ASS!!!

Obscene - you are a stupid fuck and your mother's cunt stinks

Insult - Fuck you, block me, you faggot pussy!

Toxic - What a motherfucking piece of crap those fuckheads for blocking us!

Identity - A pair of jew-hating weiner nazi schmucks.

Data preprocessing

Before applying any methods we first preprocess all of our data. We separate datasets into subgroups only, where each contains multiple documents - tweets belonging to this category. We remove retweet text RT, hyperlinks, hashtags, taggings, new lines, and zero length tweets. We further filter out tokens that not contain letters, e.g., raw punctuation.

Methodology

We start the analysis with more traditional approaches, and continue with neural approaches.

LDA

We use Latent Dirichlet Allocation (LDA) in combination with Bag-of-Words (BoW) and TF-IDF in hopes of finding obvious topics from all the provided comments / tweets. We try to determine 15 different topics, which is the same as the number of labels we have in our datasets. Results using BoW and TF-IDF are similar, however, we cannot clearly distinguish between the topics and connect obtained topics

to the existing labels, aside from one topic, which is related to sexism. Top 5 most related words are: *penis*, *rape*, *image*, *live*, *vagina*.

TF-IDF

We continue with the analysis of datasets with a traditional method TF-IDF as we want to see the most relevant words for each category of offensive language that we have in the dataset. We show the results in Table 1. We can see that some of the categories have similar unigrams that achieved the highest TF-IDF score. An example of categories with the same highest scored unigrams are *insult* and *obscene*. This makes it harder to differentiate between the categories. It is important to note, that such examples might also occur due to subjective labeling in the provided datasets, as well as people not clearly differentiating between these categories. Most datasets are not labeled by experts, but with the help of platforms such as FigureEight or Amazon Mechanical Turk. From the results in Table 1, we could assume that most people perceive categories such as *insult* and *obscene* or *threat* and *toxic* similarly. On the other hand, categories such as *spam* or *cyberbullying* are clearly differentiable from other categories. We can also see a lot of categories including Trump related words (*hatespeech*, *profane*, and *offensive*). Those categories are taken from the same dataset, and we can see that such labels will contain words that are related. So the words connected to those labels might also be connected to some bigger topic, which depends on the annotator's choice from where to extract the tweets / comments.

category	unigrams with highest TF-IDF score
racism	peopl, white, terror, man, look
sexism	feminazi, women, think, sexist, notsexist
benevolent	women, classi, sassi, nasti, gonna
abusive	know, stupid, shit, like, idiot
hateful	peopl, trump, nigga, like, idiot
spam	giveaway, game, enter, work, home
cyberbullying	one, guy, good, gone, go
hatespeech	world, trumpisatrait, trump, shameonicc, peopl
identity hate	fuck, shit, littl, like, one
insult	delet, go, ass, stupid, bitch
obscene	delet, go, stupid, bitch, ass
offensive	trumpisatrait, like, douchebag, fucktrump, get
profane	trump, shit, say, resist, peopl
threat	fuck, get, die, want, find
toxic	fuck, get, bitch, want, block

Table 1. Table shows 5 highest scoring unigrams for each label we investigate. We choose the parameters, which we believe provide us the most meaningful unigrams, so we consider words that appear in at least 5% and less than 60% of the documents.

Non-contextual word embeddings

First we select top 3 unigrams from TF-IDF results in the previous subsection (if available in the model) and find 20 most similar embeddings of pre-fitted Word2Vec ([15], [16]) and GloVe [17] (FastText unfortunately does not run due to available computational resources). We visualize the results with the help of t-SNE. We show the results in 5. From

both plots we can see that words (and its neighbors) like bitch, ass, nigga, shit, fuck, idiot, stupid, and docuhebag are relatively closely together. This could indicate a relation between *abusive*, *hateful*, *insult*, *obscene*, *identity hate*, and *toxic*. We can also see that classy, nasty, sassy, feminazi, and sexist are closely related - sometimes some words are more related in Word2Vec than in Glove and vice versa. From this we can see a relation between *sexism* and *benevolent sexism*, which is expected as they are correlated. Relationships between certain words can also vastly differ in Word2Vec and Glove. For example giveaway and game are relatively close in Word2Vec, but are further apart in Glove. Similarly, terror and white (both common unigrams of *racism* label) are far away in Word2Vec, but close in Glove. However, terror stands out from other words in both embeddings, which might indicate that racism is at least in some way different to other labels. Similarly, words from *spam* are usually more separated from words of other labels, also indicating another more clearly distinguishable group. Trump is relatively close to a lot of words in Word2Vec - women, world, peopl, sexist, die, little, guy. This could imply that general *hatespeech* and *hateful* are connected to *sexism*, *identity hate*, *threat* and also *cyberbullying*.

Additionally, after finding some initial relations with previous approaches, we also find most similar words to the category labels and compare them in a similar fashion. We show the results in Figure 6. We can see that *homophobic* and *racist* appear closely together in both embeddings and are closely related to *hateful* and *slur*. Another group which is close to the mentioned labels consist of *profane*, *vulgar*, *obscene*, *insult* and *abusive*. We can also see that *threat* and *offensive* are close together. In both embeddings *spam*, *toxic*, and *discredit* are separated from other groups and more clearly distinguishable. *Threat*, *offensive*, and *hostile* also appear closely together. From the GloVe embeddings we can also see a connection between *cyberbullying* and *harrasment*.

By now we provide some relations and decide to further investigate the connections between the related labels using word analogy. We try to find hyponyms and hypernyms, which we try with the help of the following setting:

```
father : son = our_label : x (hyponyms)
animal : cat = our_label : x (hyponyms)
son : father = our_label : x (hypernyms)
cat : animal = our_label : x (hypernyms)
```

where `our_label` is one of the analyzed labels and `x` is the word found by Word2Vec or Glove.

Unfortunately the relationships are not clear and uniquely defined. An example is *racism* is to *sexism* what is son to father with $\approx 64.6\%$ probability, but *sexism* is to *racism* what is son to father with $\approx 64.8\%$ probability. We can once again see that the two labels are related, but the precise relationship cannot be inferred. Using brother and sister the probability is lower. This could indicate that it is impossible to find a specific hypernym and that we can

only conclude that the labels are more closely related than to any other label, but they are both in some way hypernym and hyponym of each other. Similarly, *racism* and *sexism* are connected to *homophobia* and *slur*. Another group that we find, but also cannot clearly define the inner relations contain *vulgar*, *profane*, and *obscene*.

Contextual word embeddings

We move on to contextual embeddings and we focus on BERT. We use the pretrained BERT base model for Sentence Embeddings [18] and convert tweets and comments from our dataset to BERT embeddings. We use cosine similarity to compute the similarity between the embeddings. Next, we compute the average similarity between the tweets of each label to the tweets of all other labels. We show the obtained similarity matrix in Figure 7. We can see that even the tweets/comments annotated with the same label are only between 0.35 and 0.65 similar between one another according to BERT. This might indicate one of the two things - either the tweets/comments are inconsistently labeled by annotators, or BERT, even though being one of the most powerful embeddings, does not find well the contextual similarity that is specific for each subgroup of hate speech. Looking at the non-diagonal elements of the matrix, we can also see that the average similarity according to BERT between the subgroups of hate speech is around 0.5. The reason might be that the tweets/comments subgroups are very similar to the tweets/comments of the other subgroups. This might be due to the subgroups of hate speech actually being relatively similar, or again because of the annotating process.

However, for some categories we can clearly see that they are slightly less similar to other subgroups. For example *benevolent sexism* is slightly less similar to other subgroups of hate speech and it is the most similar to *cyberbullying* and *sexism*. On the other hand *cyberbullying* is more slightly more similar to most of the categories which could indicate that all of those categories are somehow included in *cyberbullying*. Another clear outlier is *spam* which is a category that is the least similar to other categories and to itself. From this we could infer that *spam* is a diverse category containing multiple different terms. We can also see that *racism* and *sexism* are less similar to other subgroups.

Discussion

Use the Discussion section to objectively evaluate your work, do not just put praise on everything you did, be critical and exposes flaws and weaknesses of your solution. You can also explain what you would do differently if you would be able to start again and what upgrades could be done on the project in the future.

References

- [1] *hate speech*. Lexico.com.




Figure 5. Word2Vec and Glove similar words. Left figure shows Word2Vec embeddings of neighboring words of top unigrams from Table 1 and the right figure shows Glove embeddings.




Figure 6. Word2Vec and Glove similar labels. Left figure shows Word2Vec embeddings of neighboring words of labels we analyze and the right figure shows Glove embeddings. Note that we omit labels that are not in vocabulary.

- [2] Ellen Spertus. Smokey: Automatic recognition of hostile messages. In *Aaai/iaai*, pages 1058–1065, 1997.
- [3] Zeerak Waseem. Are you a racist or am i seeing things? annotator influence on hate speech detection on twitter. In *Proceedings of the first workshop on NLP and computational social science*, pages 138–142, 2016.
- [4] Björn Gambäck and Utpal Kumar Sikdar. Using convolutional neural networks to classify hate-speech. In *Proceedings of the first workshop on abusive language online*, pages 85–90, 2017.
- [5] Ricardo Martins, Marco Gomes, Jose Joao Almeida, Paulo Novais, and Pedro Henriques. Hate speech classification in social media using emotional analysis. In *2018 7th Brazilian Conference on Intelligent Systems (BRACIS)*, pages 61–66. IEEE, 2018.
- [6] Marian-Andrei Rizoiu, Tianyu Wang, Gabriela Ferraro, and Hanna Suominen. Transfer learning for hate speech detection in social media. *arXiv preprint arXiv:1906.03829*, 2019.
- [7] Yi-Ling Chung, Elizaveta Kuzmenko, Serra Sinem Tekiroglu, and Marco Guerini. Conan-counter narratives through nichesourcing: a multilingual dataset of responses to fight online hate speech. *arXiv preprint arXiv:1910.03270*, 2019.
- [8] Jun-Ming Xu, Kwang-Sung Jun, Xiaojin Zhu, and Amy Bellmore. Learning from bullying traces in social media. In *Proceedings of the 2012 conference of the North American chapter of the association for computational linguistics: Human language technologies*, pages 656–666, 2012.
- [9] Carlos Arcila Calderón, Gonzalo de la Vega, and David Blanco Herrero. Topic modeling and characterization of hate speech against immigrants on twitter around the emergence of a far-right party in Spain. *Social Sciences*, 9(11):188, 2020.
- [10] Anna Schmidt and Michael Wiegand. A survey on hate speech detection using natural language processing. In *Proceedings of the fifth international workshop on natural language processing for social media*, pages 1–10, 2017.
- [11] Akshita Jha and Radhika Mamidi. When does a compliment become sexist? analysis and classification of ambivalent sexism using twitter data. In *Proceedings of the second workshop on NLP and computational social science*, pages 7–16, 2017.
- [12] Antigoni Founta, Constantinos Djouvas, Despoina Chatzakou, Ilias Leontiadis, Jeremy Blackburn, Gianluca Stringhini, Athena Vakali, Michael Sirivianos, and Nicolas Kourtellis. Large scale crowdsourcing and char-




Figure 7. Average cosine similarities between of embeddings between each label.

acterization of twitter abusive behavior. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 12, 2018.

- [13] Uwe Bretschneider and Ralf Peters. Detecting cyberbullying in online communities. 2016.
- [14] Ellery Wulczyn, Nithum Thain, and Lucas Dixon. Ex machina: Personal attacks seen at scale. In *Proceedings of the 26th International Conference on World Wide Web, WWW '17*, pages 1391–1399, Republic and Canton of Geneva, CHE, 2017. International World Wide Web Conferences Steering Committee.
- [15] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*, 2013.

- [16] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. Distributed representations of words and phrases and their compositionality. *arXiv preprint arXiv:1310.4546*, 2013.
- [17] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. Glove: Global vectors for word representation. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, 2014.
- [18] Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 11 2019.